CS 480 – Introduction to Artificial Intelligence

Topic: Adversarial Search and Games
Chapter: 5

Mustafa Bilgic

🔗 http://www.cs.iit.edu/~mbilgic

🐦 https://twitter.com/bilgicm

ADVERSARIAL SEARCH

- A **multiagent** and **competitive** environment
- Agents are trying to maximize their utilities at the expense of other agents' utilities
- Zero-sum games, where the sum of the utilities in the end is constant

  Win 1   *Draw:* $\frac{1}{2}$, $\frac{1}{2}$

  Loss 0

  - E.g., win = +1, loss = -1, tie = 0
- Games with abstract descriptions attracted the attention of AI researchers quite a bit
  - The states are relatively easy to represent, limited number of actions, precise rules, expected outcomes, etc. E.g., chess, checkers, backgammon
  - Physical games, except soccer, has not attracted much attention

UT·Austin   Peter Stone
cute robotic dogs playing Soccer

CHAPTER 5 V.S. CHAPTER 3

- Chapter 3:
  - Goal-based agent, where the path to the goal is optimized

    SA

- Chapter 5:
  - Utility-based agent, where the expected utility in the end is optimized

    MA
    competitive

Games are pretty hard to solve

- For example, chess
  - Average branching factor: ~30
  - If a game lasts 40 moves per player
    - Search tree depth is 80
    - Number of nodes: $30^{80} \approx 10^{118}$
    - Number of estimated atoms in the universe: 10^{80}
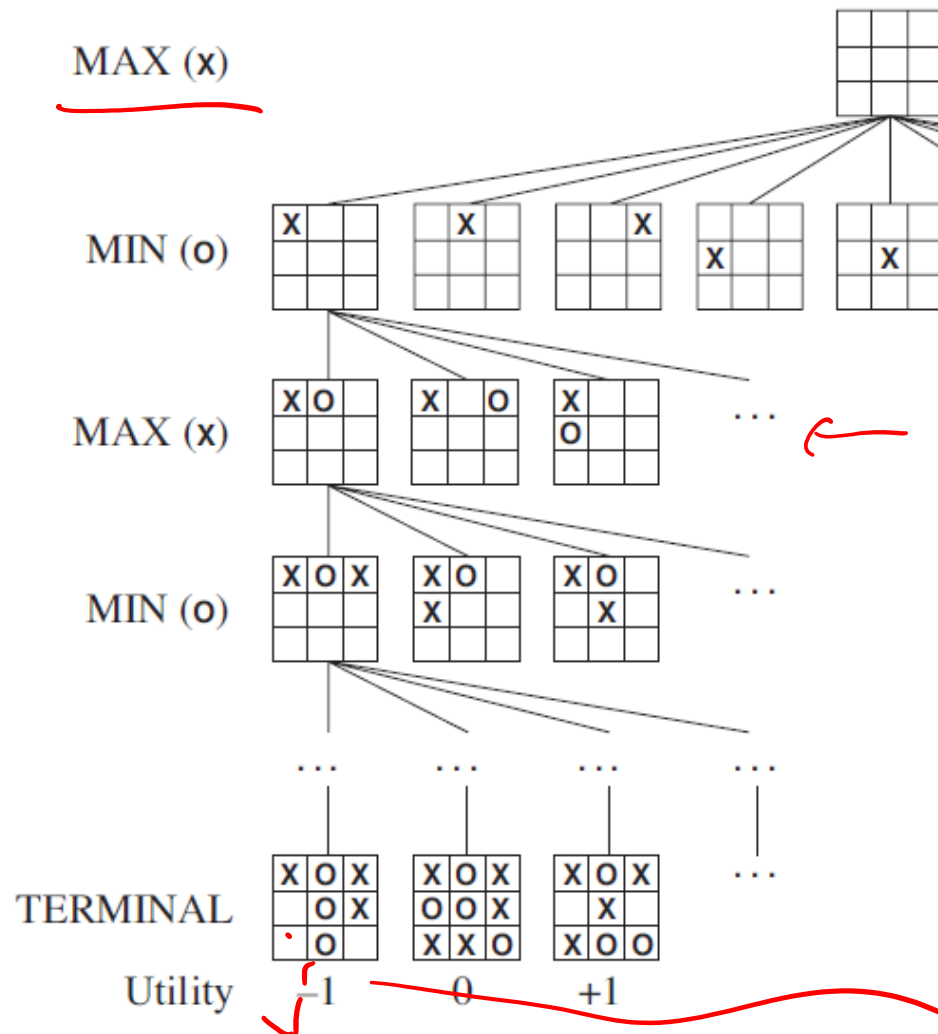 - You still must make a decision before you can calculate the optimal move

GAME DEFINITION

- S_0: The initial state

- *Player*(*s*): Which player's turn in state *s*

- *Action(s)*: The set of legal moves in state *s*

- *Result(s, a)*: The transition model

- *Terminal-test*(*s*): Is the game over

- *Utility(s, p)*: The utility of player *p* in state *s*

- Two players: MAX and MIN. MAX moves first

5

GAME TREE – TIC-TAC-TOE

$1 + 9 + 9 \cdot 8 + 9 \cdot 8 \cdot 7 + \dots$

$9!$

MAX (x)

MIN (o)   ← 9

MAX (x)   ← 9×8

What is the number of terminal nodes?

← 9×8×...

$9!$ ← 9×8×...

TERMINAL

Utility   −1   0   +1

CS 480 – Introduction to Artificial Intelligence – Illinois Institute of Technology

OPTIMAL STRATEGY

- Maximize utility

- An optimal strategy leads to outcomes at least as good as any other strategy when one is playing on *infallible* opponent

  - What if the opponent is not playing optimally?

- The **minimax** algorithm

LET'S SEE IT IN ACTION

MAX

MIN

A

a₁   a₂   a₃

B   C   D

b₁ b₂ b₃   c₁ c₂ c₃   d₁ d₂ d₃

3   12   8   2   4   6   14   5   2

TIME AND SPACE COMPLEXITY

- The minimax algorithm is a depth-first search algorithm
 - Space: $O(bm)$
  - Time: $O(b^m)$
- Can we do better?

ALPHA-BETA PRUNING

- Keep lower and upper bounds (α, β)
- MAX updates the lower bound α
- MIN updates the upper bound β
- Both pass the current bounds to their children
- MIN(α, β) prunes if $v \leq \alpha$; MAX already has a better choice α
- MAX(α, β) prunes if $v \geq \beta$; MIN already has a better choice β

EXAMPLE: MIN PRUNES

MAX $[\alpha, \beta]$

MIN $[\alpha, \beta]$

$v \nearrow$

$v < \alpha$

prune if $v \leq \alpha$

EXAMPLE: MAX PRUNES

MIN $[\alpha,\beta]$

MAX $[\alpha,\beta]$

$v \nearrow$

prune if $v \geq \beta$

$v \geq \beta$

ALPHA-BETA PRUNING ALGORITHM

function ALPHA-BETA-SEARCH($state$) **returns** an action
$v \leftarrow$ MAX-VALUE($state, -\infty, +\infty$)
return the $action$ in ACTIONS($state$) with value v

3rd

---

function MAX-VALUE($state, \alpha, \beta$) **returns** a $utility$ $value$
  if TERMINAL-TEST($state$) **then return** UTILITY($state$)
  $v \leftarrow -\infty$
  for each a **in** ACTIONS($state$) **do**
    $v \leftarrow$ MAX(v, MIN-VALUE(RESULT(s,a), α, β))
    if $v \geq \beta$ **then return** v
    $\alpha \leftarrow$ MAX(α, v)
  return v

**function** MIN-VALUE($state, \alpha, \beta$) **returns** a $utility$ $value$
  if TERMINAL-TEST($state$) **then return** UTILITY($state$)
  $v \leftarrow +\infty$
  for each a **in** ACTIONS($state$) **do**
    $v \leftarrow$ MIN(v, MAX-VALUE(RESULT(s,a), α, β))
    if $v \leq \alpha$ **then return** v
 $\beta \leftarrow$ MIN(β, v)
 **return** v

13

ALPHA-BETA PRUNING

MAX

MIN

ALPHA-BETA PRUNING – ANOTHER EXAMPLE

ALPHA-BETA PRUNING (FIGURE 5.7)

MAX

MIN

A

a_1 a_2 a_3

B   C   D

b_1 b_2 b_3 c_1 c_2 c_3 d_1 d_2 d_3

3   12   8   2   4   6   14   5   2

Can we do better?

MOVE ORDERING

MAX

MIN

What if we re-order these?

STILL

- To be able to prune, the alpha-beta algorithm has to visit some of the leaf nodes
  - Of course, this is not practical for most games
- A solution: early cut-off
  - How can we do this? Remember, the utility function is defined for the terminal states.

estimate

M

not a terminal node

state

Heuristic Evaluation Functions

- An estimate of the utility is returned rather than the actual utility

- Define features and a weighted feature combination function
  - e.g., $\text{EVAL}(s) = w_1 f_1(s) + w_2 f_2(s) + \ldots + w_n f_n(s)$

- Chess
  - Pawn: 1, Knight or bishop: 3, rook: 5, queen: 9

HEURISTIC EVALUATION FUNCTIONS

- Just like in Chapter 3, heuristic functions are not part of the problem description

  - It takes additional expertise and learning to devise useful heuristic functions

- They are not very easy to define

LINEAR COMBINATION FAILS

(a) White to move

(b) White to move

CUTTING-OFF IS NOT TRIVIAL

- A simple depth limit will not work
  - Cut-off at positions that are quiescent
    - Positions in which a drastic change is unlikely to happen
- Horizon effect
  - A serious damage is unavoidable but the search delays it by making little sacrifices
    - The little sacrifices are unnecessary as the serious damage is unavoidable but a depth-limited search cannot see it

22

Horizon Effect

Figure 5.9 The horizon effect. With Black to move, the black bishop is surely doomed. But Black can forestall that event by checking the white king with its pawns, encouraging the king to capture the pawns. This pushes the inevitable loss of the bishop over the horizon, and thus the pawn sacrifices are seen by the search algorithm as good moves rather than bad ones.

SEARCH V.S. LOOKUP

- Opening of a game
  - A depth-limited search is unlikely to help
  - Human expertise and learning from previous games for opening moves

- Ending a game
  - When very few pieces are left, the number of possible states is not many; rather than a depth-limited search, look-up a winning strategy

Monte-Carlo Tree Search

function Monte-Carlo-Tree-Search(*state*) **returns** *an action*
    tree ← Node(*state*)
    while Is-Time-Remaining() **do**
        leaf ← Select(*tree*)
        child ← Expand(*leaf*)
        result ← Simulate(*child*)
        Back-Propagate(*result*, *child*)
    return the move in Actions(*state*) whose node has highest number of playouts

Figure 5.11 The Monte Carlo tree search algorithm. A game tree, *tree*, is initialized, and then we repeat a cycle of Select / Expand / Simulate / Back-Propagate until we run out of time, and return the move that led to the node with the highest number of playouts.

MONTE-CARLO TREE SEARCH

(a) Selection

(b) Expansion and simulation

black wins

(c) Backpropagation

Figure 5.10 One iteration of the process of choosing a move with Monte Carlo tree search (MCTS) using the upper confidence bounds applied to trees (UCT) selection metric, shown after 100 iterations have already been done. In (a) we select moves, all the way down the tree, ending at the leaf node marked 27/35 (for 27 wins for black out of 35 playouts). In (b) we expand the selected node and do a simulation (playout), which ends in a win for black. In (c), the results of the simulation are back-propagated up the tree.

So Far

- Deterministic
- Fully-observable

Stochastic Games

- An element of chance, such as dice

- Modify the search tree and the minimax algorithm:

  - Include a chance node, that leads to all possible outcomes

  - Rather than computing utilities, compute expected utilities

28

PARTIALLY OBSERVABLE GAMES

- The current state is only partially observable

- Model what is not observed

- Examples
  - Card games
  - Battleships
  - Kriegspiel

29

HISTORICAL NOTES

- Algorithms
  - 1912 – Minimax
  - 1956-1961 – Alpha-beta pruning
  - 1987 – Monte Carlo tree search
- Games
  - Chess
    - 1982 – Master status
    - 1990 – International master status
    - 1997 – Defeated Gary Kasparov, the world champion
    - 2017 – AlphaZero defeated Stockfish
  - Checkers
    - 1960s – learned from self-play using reinforcement learning
    - 2007 – The game is fully solved

30

HISTORICAL NOTES

- Games (continued)
  - Go
    - 1970 – Visual pattern recognition for Go was proposed
    - 1993 – Monte-Carlo tree search for Go
    - 1994 – reinforcement learning for Go
    - 2001 – neural networks for Go
    - 2016 – AlphaGo put all these ideas together
  - Backgammon
    - 1980s – first computer programs for Backgammon
    - 1995 – TD-Gammon – neural networks and self-play – able to play at world champion level
  - Poker
    - 2015 – Optimal strategy was developed (using game theory)
    - 2017 – Computers beat world champions
  - StarCraft II
    - 2019 – AlphaStar defeats expert players
      - Discussion whether programs have unfair advantage in action speed